



CS 329P : Practical Machine Learning (2021 Fall)

3.3 Linear Methods

Qingqing Huang, Mu Li, Alex Smola

<https://c.d2l.ai/stanford-cs329p>

Linear Regression



- A simple house price prediction
 - Assume 3 features: $x_1 = \text{\#beds}$, $x_2 = \text{\#baths}$, $x_3 = \text{living sqft}$
 - The predicted price is $\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b$
 - Weights w_1, w_2, w_3 and bias b will be learnt from training data
- In general, given data $\mathbf{x} = [x_1, x_2, \dots, x_p]$, linear regression predicts

$$\hat{y} = w_1x_1 + w_2x_2 + \dots + w_px_p + b = \langle \mathbf{w}, \mathbf{x} \rangle + b$$

- Code:

```
# weight w has shape (p, 1)
# bias b is a scalar
# data x has shape (p, 1)
y_hat = (x*w).sum() + b
```

Objective Function



- Collect n training examples $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_1, \dots, \mathbf{x}_n]^T \in \mathbb{R}^{n \times p}$ with labels $\mathbf{y} = [y_1, \dots, y_n]^T \in \mathbb{R}^n$
 - E.g. house listings and sold prices
- Objective: minimize the mean square error (MSE)


$$\begin{aligned}\mathbf{w}^*, \mathbf{b}^* &= \operatorname{argmin}_{\mathbf{w}, b} \ell(\mathbf{X}, \mathbf{y}, \mathbf{w}, b) \\ &= \operatorname{argmin}_{\mathbf{w}, b} \frac{1}{n} \sum_{i=1}^n (y_i - \langle \mathbf{x}_i, \mathbf{w} \rangle - b)^2\end{aligned}$$

- Exercise: write the closed form solution

Use linear regression for classification problem



- Regression: continuous output in \mathbb{R}
- Multi-class classification:
 - One-hot label $\mathbf{y} = [y_1, y_2, \dots, y_m]$ where $y_i = 1$ if $i = y$ otherwise 0
 - $\hat{\mathbf{y}} = \mathbf{o}$, where i -th output o_i is the confidence score for class i
 - Learn a liner model for each class $o_i = \langle \mathbf{x}, \mathbf{w}_i \rangle + b_i$
 - Minimize MSE loss $\frac{1}{m} \|\mathbf{o} - \mathbf{y}\|_2^2$
 - Predict label $\operatorname{argmax}_i \{o_i\}_{i=1}^m$
- Waste model capacity on pushing o_i near 0 for off labels



Softmax Regression




- One-hot label $\mathbf{y} = [y_1, y_2, \dots, y_m]$ where $y_i = 1$ if $i = y$ otherwise 0

- $\hat{\mathbf{y}} = \text{softmax}(\mathbf{o})$ where $\hat{y}_i = \frac{\exp(o_i)}{\sum_{k=1}^m \exp(o_k)}$

- Turns confidence scores into probabilities (non-negative, sum to 1)
- Ideally we want $\hat{\mathbf{y}} = \text{one-hot}(\text{argmax}_i o_i)$, softmax is a continuous approximate to that
- Still a linear model, decision made on linear transformation of the input, as $\text{argmax}_i \hat{y}_i = \text{argmax}_i o_i$
- Cross-entropy loss between two distributions $\hat{\mathbf{y}}$ and \mathbf{y}

$$H(\mathbf{y}, \hat{\mathbf{y}}) = \sum_i -y_i \log(\hat{y}_i) = -\log \hat{y}_y$$

- When label class is i , assigns less penalty on o_j as long as $o_j \ll o_i$
- Exercise: think about how to handle examples with multi labels?



Mini-batch Stochastic gradient descent (SGD)




- Train by mini-batch SGD (by various other ways as well)
 - \mathbf{w} model param, b batch size, η_t learning rate at time t
 - Randomly initialize \mathbf{w}_1
 - Repeat $t = 1, 2, \dots$ until converge
 - Randomly samples $I_t \subset \{1, \dots, n\}$ with $|I_t| = b$
 - Update $\mathbf{w}_{t+1} = \mathbf{w}_t - \eta_t \nabla_{\mathbf{w}_t} \ell(\mathbf{X}_{I_t}, \mathbf{y}_{I_t}, \mathbf{w}_t)$
- Pros: solve all objectives in this course except for trees
- Cons: sensitive to hyper-parameters b and η_t



Code

- Train a linear regression model with min-batch SGD
- Hyperparameters
 - batch_size
 - learning_rate
 - num_epochs

Full code at: http://d2l.ai/chapter_linear-networks/linear-regression-scratch.html



```
# `features` shape is (n, p), `labels` shape is (n, 1)
def data_iter(batch_size, features, labels):
    num_examples = len(features)
    indices = list(range(num_examples))
    random.shuffle(indices) # read examples at random
    for i in range(0, num_examples, batch_size):
        batch_indices = torch.tensor(
            indices[i:min(i + batch_size, num_examples)])
        yield features[batch_indices], labels[batch_indices]

w = torch.normal(0, 0.01, size=(p, 1), requires_grad=True)
b = torch.zeros(1, requires_grad=True)

for epoch in range(num_epochs):
    for X, y in data_iter(batch_size, features, labels):
        y_hat = X @ w + b
        loss = ((y_hat - y)**2 / 2).mean()
        loss.backward()
        for param in [w, b]:
            param -= learning_rate * param.grad
            param.grad.zero_()
```

Summary



- Linear methods linearly combine inputs to obtain predictions
- Linear regression uses MSE as the loss function
- Softmax regression is used for multiclass classification
 - Turn predictions into probabilities and use cross-entropy as loss
 - Cross entropy loss between two probability distribution
- Mini-batch SGD can learn both models (and later neural networks as well)