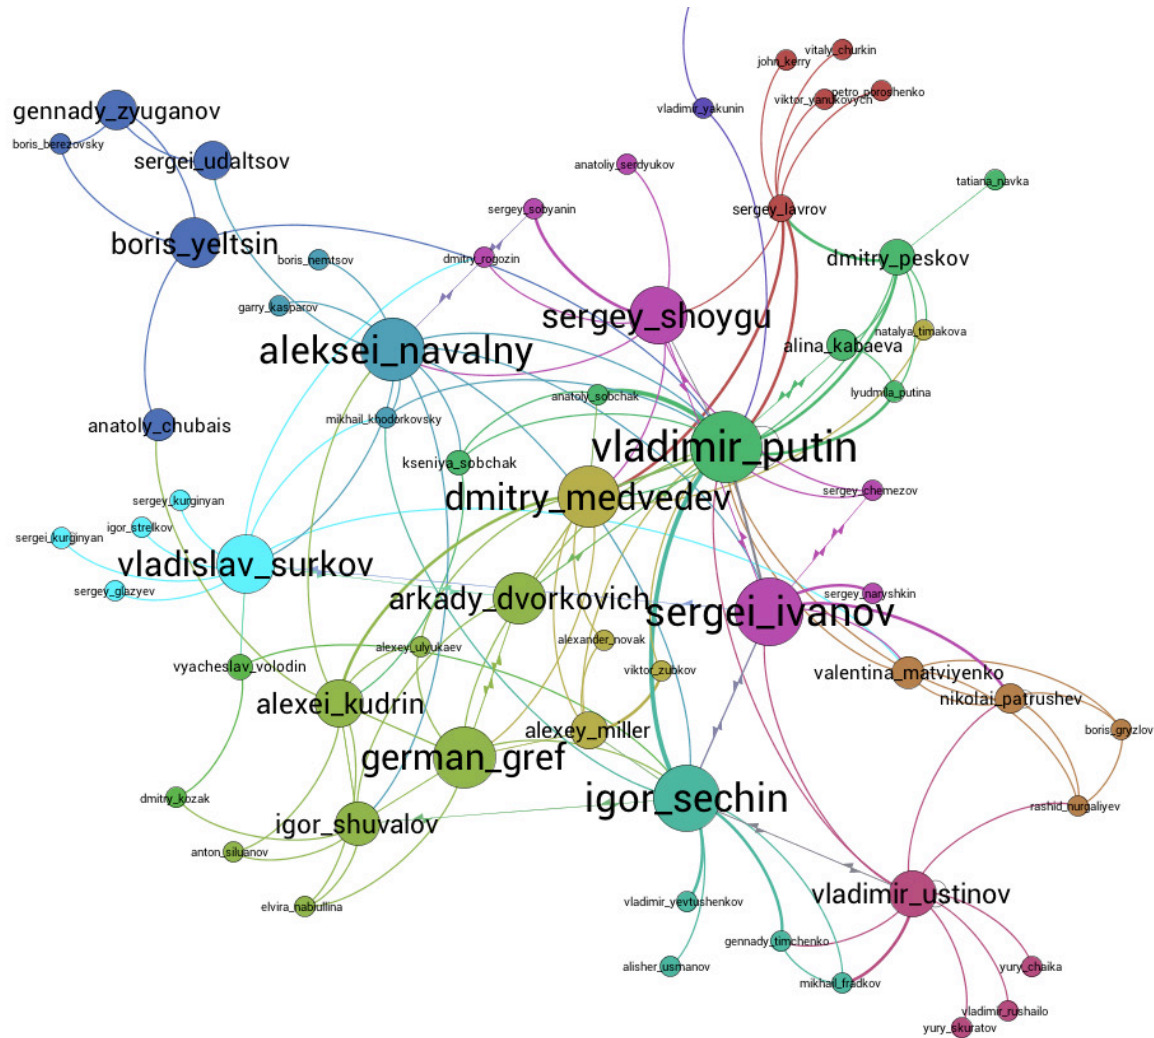


Graphs



Beyond Sequential Dependency



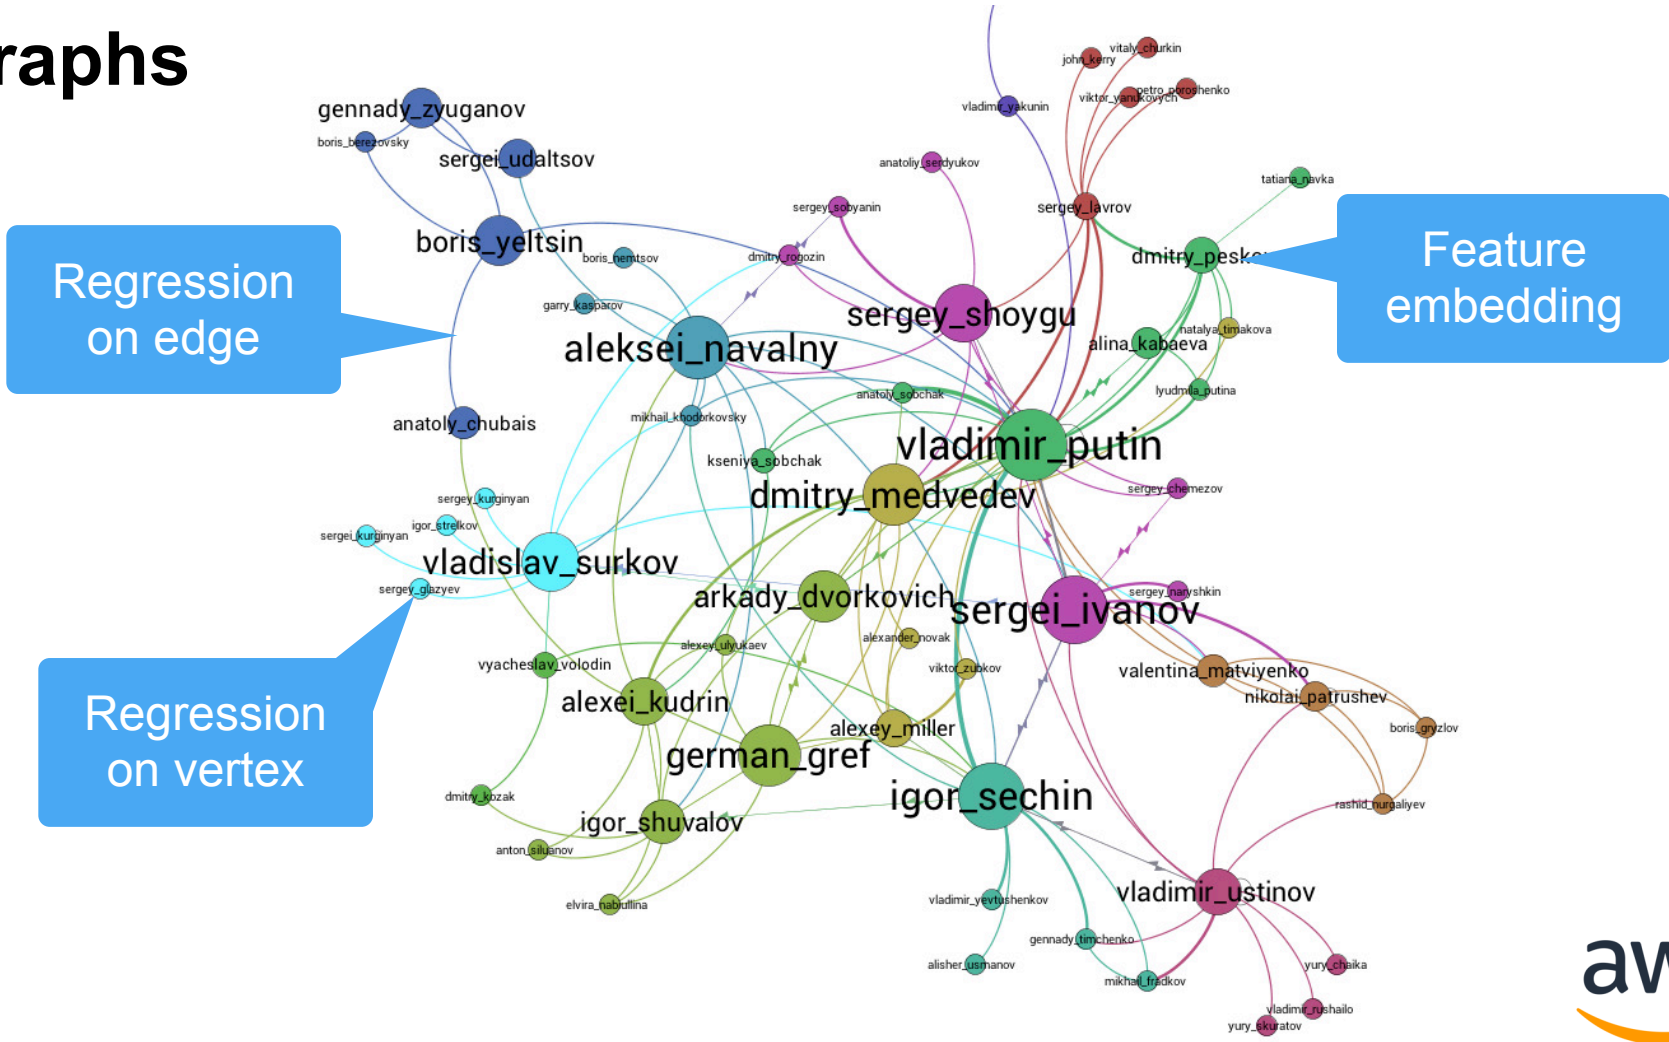
Often difficult to model

- **Spatial dependency** (e.g., road network)
- **General graph dependency**
 - Spread of memes, fake news ...
 - Product recommendation (users, items, vendors)
 - Relational databases
- Cannot write as sequence but needs graph. Popular choices:

Directed graphical model $p(x) = \prod_i p(x_i | x_{\pi(i)})$

Undirected graphical model $p(x) = \prod_C \psi_C(x_C)$

Graphs

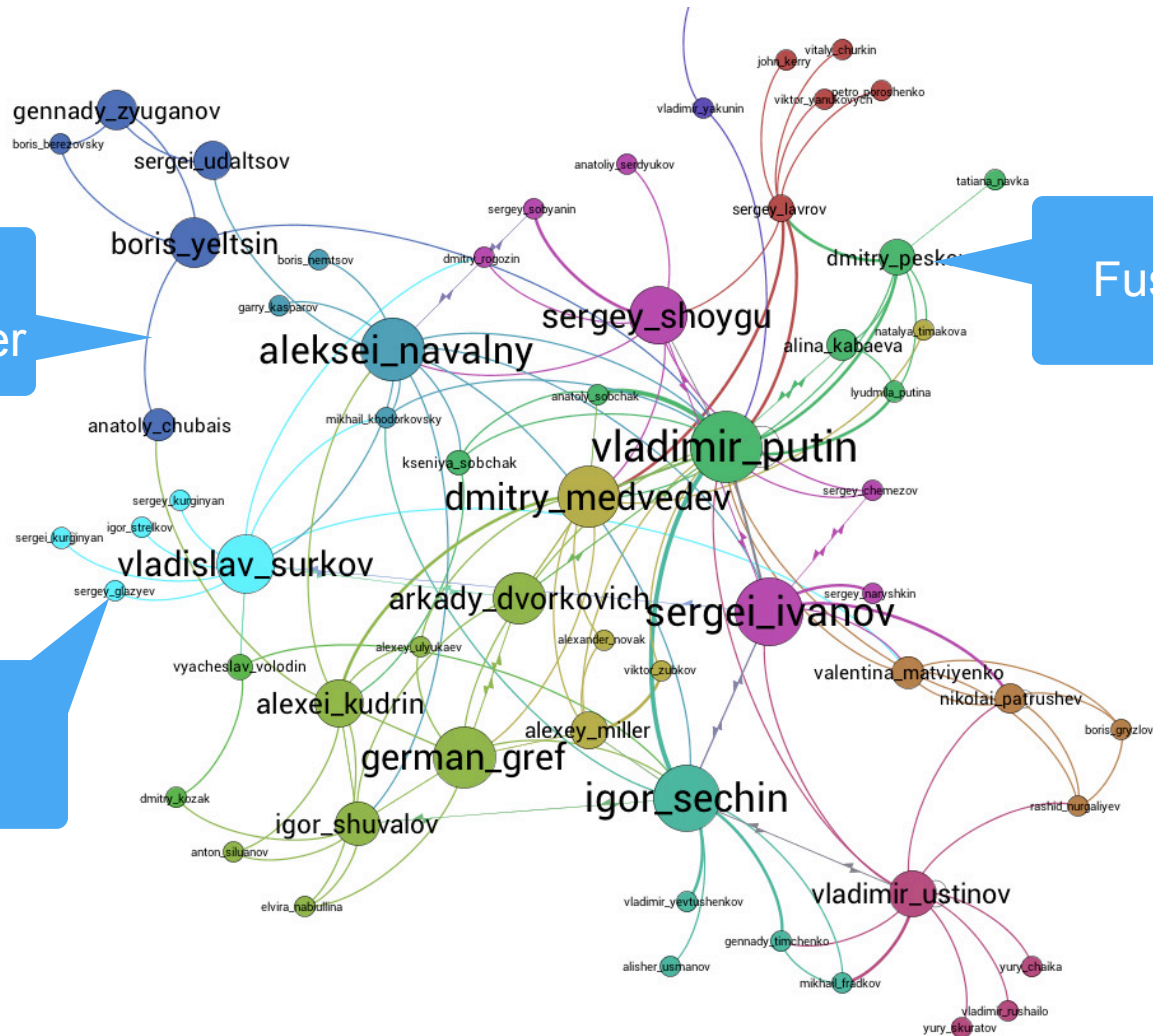


Graphs

Social
Recommender

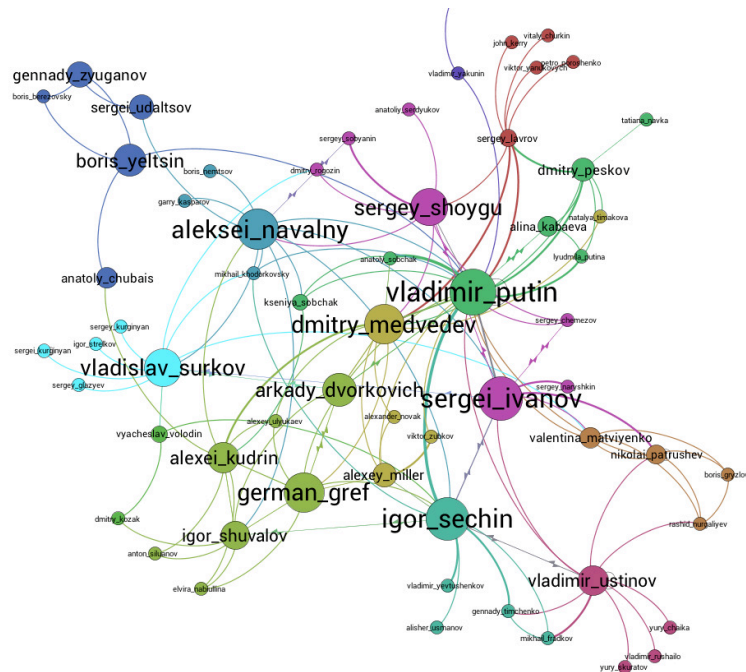
Fraud
detection

Fuse data



Graphs $G(V,E)$

- Vertices $i \in V$ (with attributes)
- Edges $(i,j) \in E$ (with attributes)
- Estimation problems
 - Given some vertex labels y estimate the remaining ones (e.g. fraud detection)
 - Given some edge attributes e estimate the remaining ones (e.g. link recommendation)



Vertex Updates



Weisfeiler-Lehman algorithm (1976)

- **Key idea**

- Graph isomorphism ... is trivial if vertices are unique
- Make them unique by repeated hashing

$$v(i) \leftarrow h(v(i), \{v(j) \text{ with } j \sim i\})$$

Local updates

- Terminate once stationary
- **Machine Learning variant** (Shervashidze & Borgwardt, 2013)
 - The vertex hashes are good features $i \rightarrow \phi(i)$
 - Can prove equivalence to some graph kernels
- Crazy thought ... what about vertices with attributes?

Page Rank (Page, Brin, Motwani, Kleinberg, 1990s)

- **Random surfer model with restarts**

- Start at uniformly random location
- Follow random link at each vertex
- Page rank is stationary distribution

- **Self consistency equation**

$$p(i) = \frac{\epsilon}{|V|} + (1 - \epsilon) \sum_{i \sim j} \frac{p(j)}{d(j)} \iff r(i) = \epsilon + (1 - \epsilon) \sum_{i \sim j} \frac{r(j)}{d(j)}$$

- Solve by iterating fixed number of times or until convergence

Random
Restart

Random
Link following

Local updates

(naive) PageRank Implementation in DGL

```
import dgl.function as fn
```

```
def pagerank_builtin(g):  
    g.ndata['pv'] = g.ndata['pv'] / g.ndata['deg']  
    g.update_all(message_func=fn.copy_src(src='pv', out='m'),  
                  reduce_func=fn.sum(msg='m', out='m_sum'))  
    g.ndata['pv'] = (1 - DAMP) / N + DAMP * g.ndata['m_sum']
```

Incoming PageRank

Random surfer

Random surfer

Belief Propagation / Message Passing

- **Graphical models** (we operate on clique graph)
- Vertex potential and directed messages

$$\phi_C = \psi_C \prod_{D \sim C} \mu_{D \rightarrow C}$$

$$\mu_{D \rightarrow C}(x_{C \cap D}) = \sum_{x_{C \setminus D}} \phi_C(x_C) \mu_{D \rightarrow C}^{-1}(x_{C \cap D})$$

- Finite time convergence if clique graph is a junction tree
- In practice ignore and run *for some time* on clique graph (loopy belief propagation)

Belief Propagation / Message Passing

- **Graphical models** (we operate on clique graph)
- Vertex potential and directed messages

$$\phi_C = \psi_C \prod_{D \sim C} \mu_{D \rightarrow C}$$

$$\mu_{D \rightarrow C}(x_{C \cap D}) = \sum_{x_{C \setminus D}} \phi_C(x_C) \mu_{D \rightarrow C}^{-1}(x_{C \cap D})$$

- Crazy thought ... what if we didn't care about graphical models (they might not converge anyway)?

This looks like vertex updates
Can we learn them?
Can we get vertex features?

Graph Convolutions (e.g. Kipf & Welling, 2016)

- **Basic idea**

- Vertex (and edge) features x_i and x_{ij}
- Compute new vertex (and edge) features using local update function

$$v_i \leftarrow f(v_i, \{v_j, v_{ij} \text{ with } i \sim j\})$$

$$v_{ij} \leftarrow g(v_{ij}, v_i, v_j)$$

Optional (e.g. not in K&W'16)

- **Variants**

- Run for a fixed number of steps (like graph kernel)
- Run to convergence (like page rank)

Forward model

```
G.ndata['feat'] = torch.eye(34)
```

```
def gcn_message(edges):  
    # The argument is a batch of edges using source node's feature 'h'  
    return {'msg' : edges.src['h']}
```

```
def gcn_reduce(nodes):  
    # The argument is a batch of nodes with features summed over 'msg'  
    return {'h' : torch.sum(nodes.mailbox['msg'], dim=1)}
```

```
class GCNLayer(nn.Module):  
    def __init__(self, in_feats, out_feats):  
        super(GCNLayer, self).__init__()  
        self.linear = nn.Linear(in_feats, out_feats)  
  
    def forward(self, g, inputs):  
        g.ndata['h'] = inputs  
        g.send(g.edges(), gcn_message) # trigger message passing  
        g.recv(g.nodes(), gcn_reduce) # trigger aggregation at all nodes  
        h = g.ndata.pop('h') # get the result node features  
        return self.linear(h) # perform linear transformation
```



Forward Model

```
class GCN(nn.Module):
    def __init__(self, in_feats, hidden_size, num_classes):
        super(GCN, self).__init__()
        self.gcn1 = GCNLayer(in_feats, hidden_size)
        self.gcn2 = GCNLayer(hidden_size, num_classes)

    def forward(self, g, inputs):
        h = self.gcn1(g, inputs)
        h = torch.relu(h)
        h = self.gcn2(g, h)
        return h

# First layer - 34 inputs to 5 dimensions
# Second layer - 2 dimensional embedding for 2 groups
net = GCN(34, 5, 2)
```

Training it

```
optimizer = torch.optim.Adam(net.parameters(), lr=0.01)
all_logits = []
for epoch in range(30):
    logits = net(G, inputs)
    # save the logits for visualization later
    all_logits.append(logits.detach())
    logp = F.log_softmax(logits, 1)
    # only compute loss for labeled nodes
    loss = F.nll_loss(logp[labeled_nodes], labels)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

Graph Convolutions with state (GeniePath - Liu et al., 2018, Platanios & S, 2018)

- **Basic idea**

- Vertex (and edge) features x_i and x_{ij}
- Vertex has internal state, LSTM or similar (GeniePath)
- Compute new vertex (and edge) features using local update function

$$(v_i^{t+1}, h_i^{t+1}, c_i^{t+1}) \leftarrow \text{LSTM}(v_i^t, h_i^t, c_i^t, \{v_j, v_{ij} \text{ with } i \sim j\})$$

$$v_{ij} \leftarrow g(v_{ij}, v_i, v_j)$$

Optional

- This works better on some datasets

Graph Convolutions with state (GeniePath - Liu et al., 2018, Platanios & S, 2018)

- **Basic idea**

- Vertex (and edge) features x_i and x_{ij}
- Vertex has internal state, LSTM or similar (GeniePath)

- **Minor twist (Deep Sets - Zaheer et al., 2017)**

All functions on sets are nonlinear sums

$$(v_i^{t+1}, h_i^{t+1}, c_i^{t+1}) \leftarrow \text{LSTM} \left(v_i^t, h_i^t, c_i^t, g \left(\sum_{j \sim i} f(v_j, v_{ij}) \right) \right)$$

$$v_{ij} \leftarrow g(v_{ij}, v_i, v_j)$$

Graph Attention Networks (Velickovic et al., 2017)

- **Basic idea**

- Aggregation from neighboring vertices should be attention gated (self attention strategy)

- Without attention

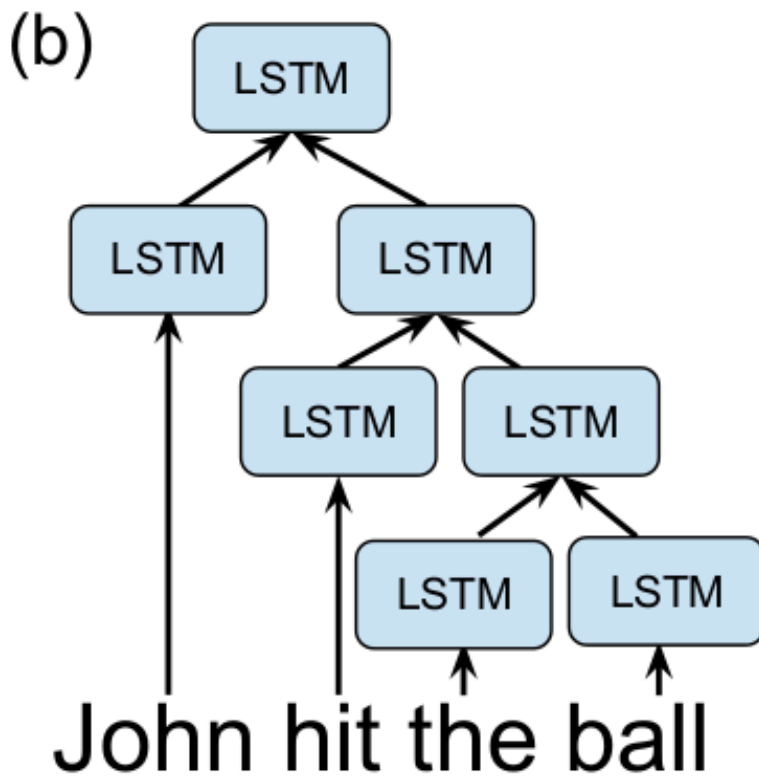
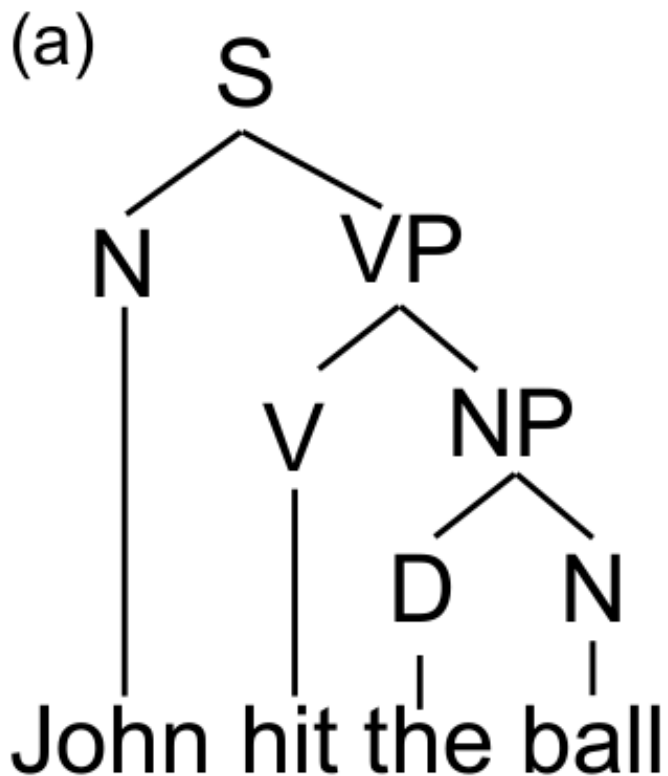
$$v_i \leftarrow f\left(v_i, \sum_{i \sim j} h(v_j)\right)$$

- With attention

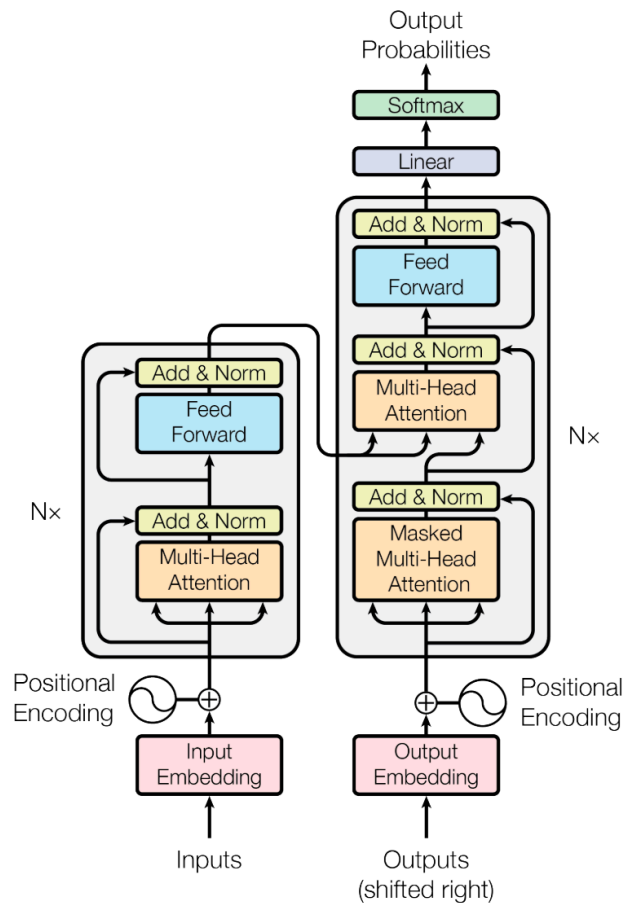
$$v_i \leftarrow f\left(v_i, \sum_{i \sim j} \alpha_{ij} h(v_j)\right) \text{ where } \alpha_{ij} = \text{softmax}\left(\{v_i^\top M v_j\}\right)$$

- This works better on some datasets

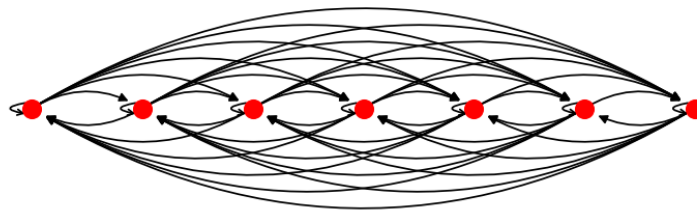
Tree-LSTM as message-passing



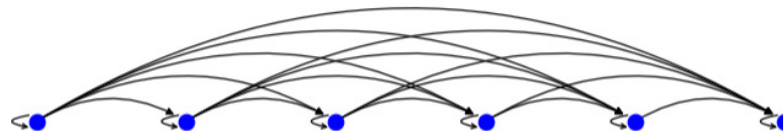
Deconstruct Transformer as a graph



Encoder

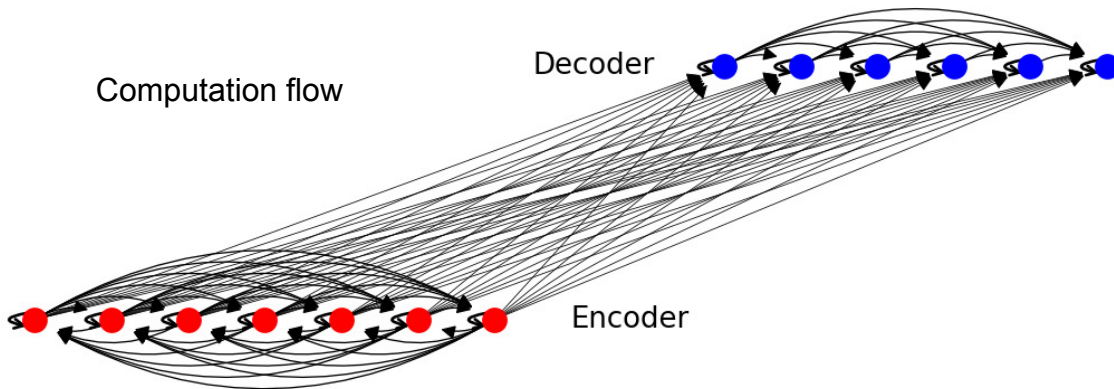


Decoder



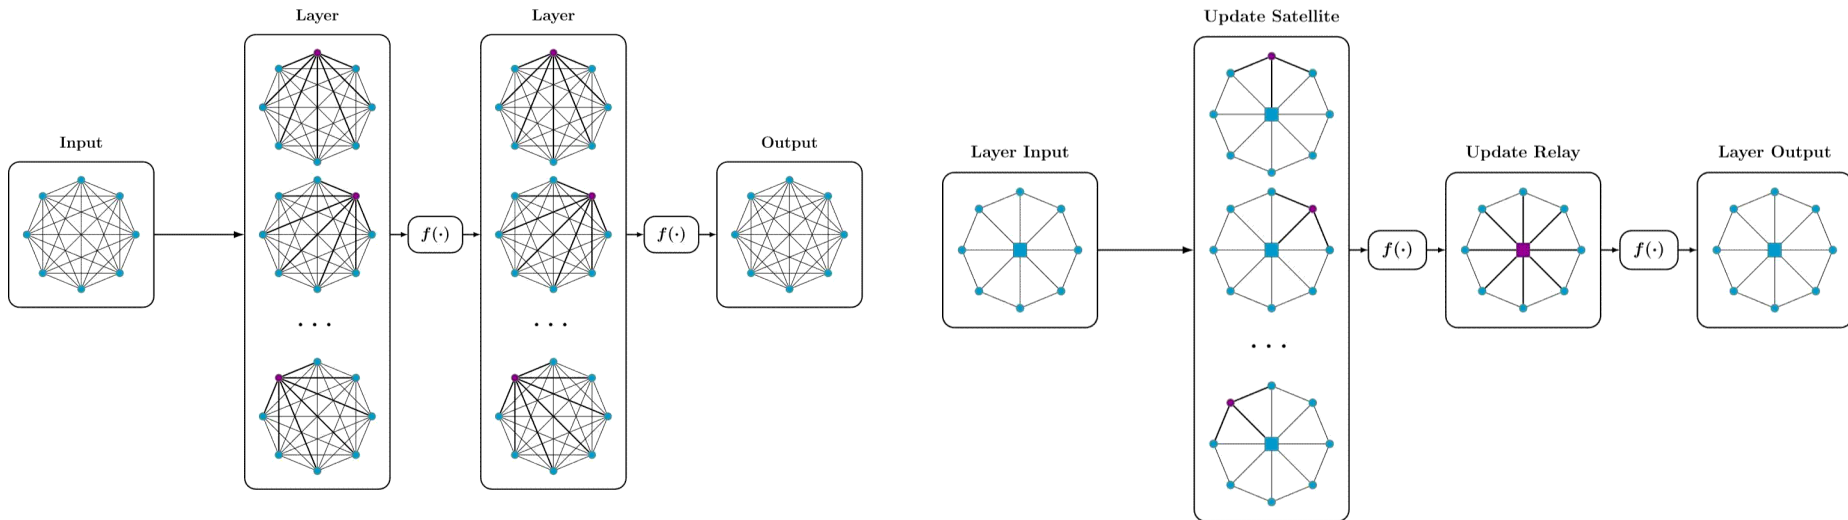
Computation flow

Decoder



Encoder

Deconstruct Transformer as a graph



Transformer:

- data & compute hungry

[Star-Transformer](#) (in NAACL'19)

- much less data hungry
- leverages ngram prior,
- has issues with long dependency

[SegTree-Transformer](#) (in ICLR'19 RLGM)

- less data hungry
- A good compromise in between

Making it Practical (Dai et al., 2018)

- Learning the vertex update function is **expensive**
 - Backprop over graph has to deal with entire graph quickly (6 degrees of separation kill BPTT)
 - Not much benefit in finite iterations
- Replace with fixed point iteration
 - Can learn it directly
 - Local convergence
 - **Sample vertex updates** (much smaller subset)
- This works better on some datasets

Making it Practical (Dai et al., 2018)

- Initialize
- Compute update
$$v_i \leftarrow f(v_i, \{v_j \text{ with } i \sim j\})$$
- Compute (regression) loss from embedding
- Backprop to change
 - f via loss
 - f via self consistency

