



CS 329P : Practical Machine Learning (2021 Fall)

9.3 NAS algorithms

Qingqing Huang, Mu Li, Alex Smola

<https://c.d2l.ai/stanford-cs329p>

Neural Architecture Search (NAS)



- A neural network has different types of hyperparameters:

- Topological structure: resnet-ish, mobilenet-ish, #layers
- Individual layers: kernel_size, #channels in convolutional layer, #hidden_outputs in dense/recurrent layers



- NAS automates the design of neural network

- How to specify the search space of NN
- How to explore the search space
- Performance estimation

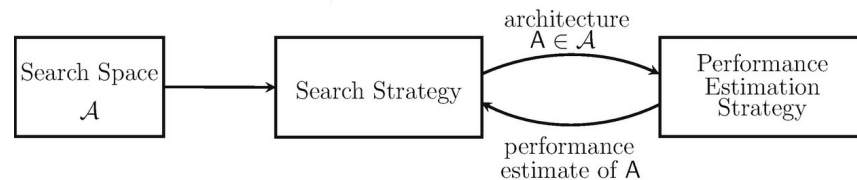


Image source: [Elsken, et al. 2019](#)

NAS with Reinforcement Learning



- Zoph & Le 2017



- A RL-based controller (REINFORCE) for proposing architecture.
 - RNN controller outputs a sequence of tokens to config the model architecture.
 - Reward is the accuracy of a sampled model at convergence
- Naive approach is expensive and sample inefficient (~2000 GPU days). To speed up NAS:
 - Estimate performance
 - Parameter sharing (e.g. EAS, ENAS)

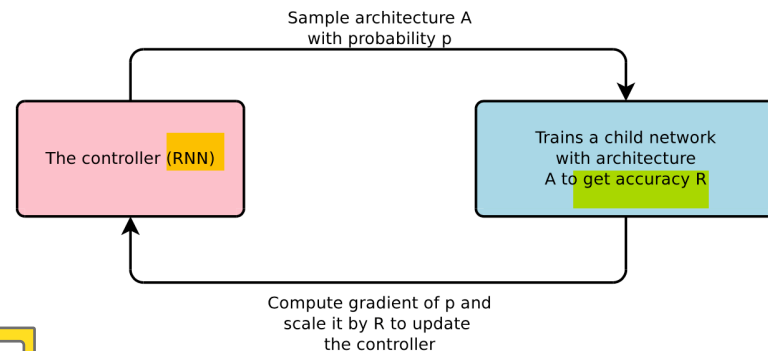



Image source: [Zoph & Le 2017](#)

The One-shot Approach

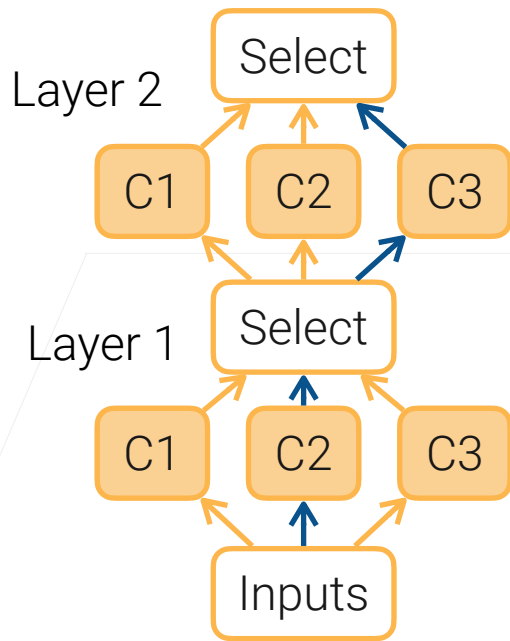


- Combines the learning of architecture and model params
- Construct and train a single model presents a wide variety of architectures
- Evaluate candidate architectures 
 - Only care about the candidate ranking
 - Use a proxy metric: the accuracy after a few epochs
- Re-train the most promising candidate from scratch

Differentiable Architecture Search



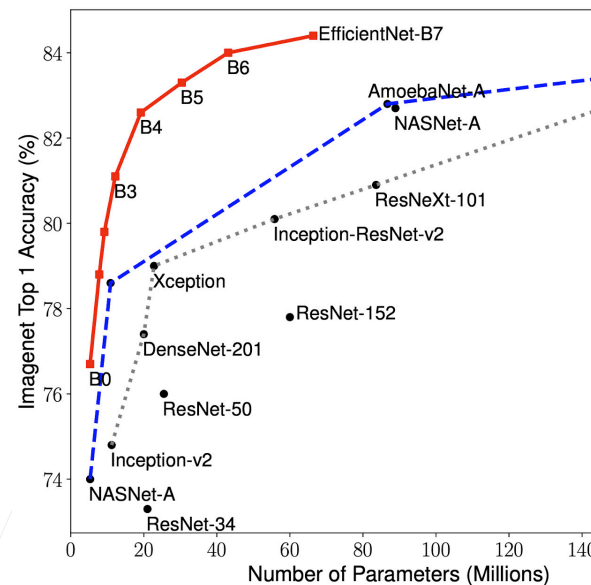
- Relax the categorical choice to a softmax over possible operations:
 - Multiple candidates for each layer
 - Output of i -th candidate at layer l is o_i^l
 - Learn mixing weights \mathbf{a}^l . The input for $i + 1$ -the layer is $\sum_i \alpha_i^l o_i^l$ with $\alpha^l = \text{softmax}(\mathbf{a}^l)$
 - Choose candidate $\text{argmax}_i \alpha_i$
 - Jointly learn \mathbf{a}^l and network parameters
- A more sophisticated version (DARTS) achieves SOTA and reduces the search time to ~3 GPU days



Scaling CNNs







- A CNN can be scaled by 3 ways:
 - Deeper: more layers
 - Wider: more output channels
 - Larger inputs: increase input image resolutions
- EfficientNet proposes a compound scaling
 - Scale depth by α^ϕ , width by β^ϕ , resolution by γ^ϕ
 - $\alpha\beta^2\gamma^2 \approx 2$ so increase FLOP by 2x if $\phi = 1$
 - Tune $\alpha, \beta, \gamma, \phi$



Research directions



- Explainability of NAS result
- Search architecture to fit into edge devices 
 - Edge devices are more and more powerful, data privacy concerns
 - But they are very diverse (CPU/GPU/DSP, 100x performance difference) and have power constraints 
 - Minimize both model loss and hardware latency 
 - E.g. minimize loss $\times \log(\text{latency})^\beta$
- To what extent can we automate the entire ML workflow? 

Summary



- NAS searches a NN architecture for a customizable goal
 - Maximize accuracy or meet latency constraints on particular hardware
- NAS is practical to use now:
 - Compound depth, width, resolution scaling
 - Differentiable one-hot neural network