



CS 329P : Practical Machine Learning (2021 Fall)

11. Transfer Learning

Qingqing Huang, Mu Li, Alex Smola

<https://c.d2l.ai/stanford-cs329p>

Transfer learning



- Motivation



- Exploit a model trained on one task for a related task
- Popular in deep learning as DNNs are data hungry and training cost is high

- Approaches

- Feature extraction (e.g. Word2Vec, ResNet-50 feature, 3D feature)
- Train a model on a related task and reuse it
- Fine-tuning from a pretrained model (focus of this lecture)



- Related to

- Semi-supervised learning
- In the extreme, zero-shot / few-shot learning
- Multi-task learning, where some labeled data is available for each task





CS 329P : Practical Machine Learning (2021 Fall)



11.1 Fine-tuning in CV

Qingqing Huang, Mu Li, Alex Smola




<https://c.d2l.ai/stanford-cs329p>

Transferring Knowledge



- There exists **large-scale labeled CV** datasets 
- Especially for **image classification**, the cheapest one to label
- Transfer knowledge from models trained on these datasets to your CV applications (with 10-100X smaller data) 



Your dataset




# examples	1.2 M	50K	60 K
# classes	1,000	100	10

Pre-trained Models



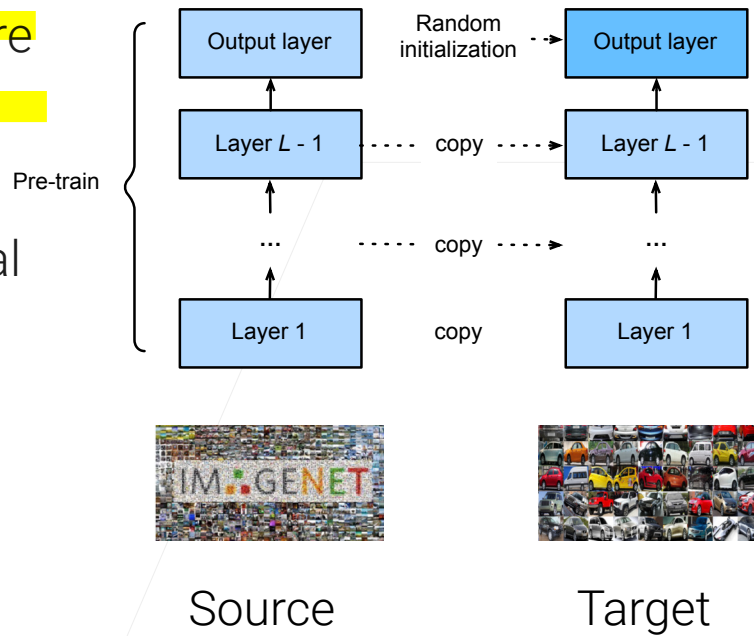
- Partition a neural network into: 
 - A feature extractor (**encoder**) maps raw pixels into linearly separable features
 - A linear classifier (**decoder**) makes decisions
- **Pre-trained model** 
 - a neural network trained on a large-scale and general enough dataset
 - The feature extractor may **generalize well** to
 - **other datasets** (e.g. medical/satellite images)
 - **other tasks** (e.g. object detection, segmentation)



Fine-Tuning techniques



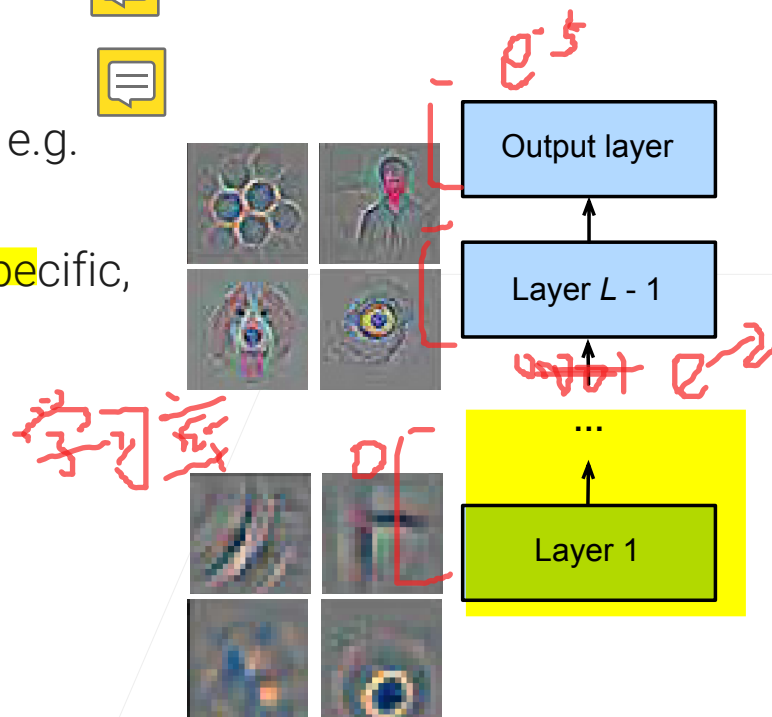
- Initialize the new model:
 - Initialize the feature extractor with the feature extractor parameters of a pre-trained model
 - Randomly initialize the output layer
 - Start the parameter optimization near a local minimal
- Train with a small learning rate with just a few epochs
 - Regularize the search space



Freeze Bottom Layers



- Neural networks learn hierarchical features
 - Low-level features are universal, generalize well, e.g. curves / edges / blobs
 - High-level features are more task and dataset specific, e.g. classification labels
- Freeze bottom layers during fine tuning
Train the top layers from scratch
 - Keep low-level universal features intact
 - Focus on learning task specific features
 - A strong regularizer



Where to Find Pre-trained Models





- Tensorflow Hub: <https://tfhub.dev/>
 - Tensorflow models submitted by users
- TIMM: <https://github.com/rwightman/pytorch-image-models>
 - PyTorch models collected by Ross Wightman



```
import timm
from torch import nn
```

```
model = timm.create_model('resnet18', pretrained=True)
model.fc = nn.Linear(model.fc.in_features, n_classes)
# Train model as a normal training job
```


Applications



- Fine-tuning pre-trained models (on ImageNet) is widely used in various CV applications: 
 - Detection/segmentation (similar images but different targets)
 - Medical/satellite images (same task but very different images)
- Fine-tuning accelerates convergence 
- Though not always improve accuracy
 - Training from scratch could get a similar accuracy, especially when the target dataset is also large



Summary



- Pre-train models on large-scale datasets (often image classification)
- Initialize weights with pre-trained models for down-stream tasks
- Fine-tuning accelerates converges and (sometimes) improves accuracy